

Business Intelligence: OLAP, Data Warehouse, and Column Store

Thomas Heinis

Why we still study OLAP/Data Warehouse in Big Data?

- Understand the Big Data history
 - How does the requirement of (big) data analytics/business intelligence evolve over the time?
 - What are the architecture and implementation techniques being developed? Will they still be useful in Big Data?
 - Understand their limitation and what factors have changed from 90's to now?
- NoSQL is not only SQL 😊
- Hive/Impala aims to provide OLAP/BI for Big Data using Hadoop

Highlights

- OLAP
 - Multi-relational Data model
 - Operators
 - SQL
- Data warehouse (architecture, issues, optimizations)
- Join Processing
- Column Stores (Optimized for OLAP workload)

Back to the 70's: Relational Databases

Basic Structure

- Formally, given sets D_1, D_2, \dots, D_n a **relation** r is a subset of $D_1 \times D_2 \times \dots \times D_n$

Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$

- Example:

customer_name = {Jones, Smith, Curry, Lindsay}

customer_street = {Main, North, Park}

customer_city = {Harrison, Rye, Pittsfield}

Then $r = \{$ (Jones, Main, Harrison),
 (Smith, North, Rye),
 (Curry, North, Rye),
 (Lindsay, Park, Pittsfield) $\}$

is a relation over

customer_name, *customer_street*, *customer_city*

Relation Schema

- A_1, A_2, \dots, A_n are *attributes*
- $R = (A_1, A_2, \dots, A_n)$ is a *relation schema*

Example:

Customer_schema = (customer_name, customer_street, customer_city)

- $r(R)$ is a *relation* on the *relation schema* R

Example:

customer (Customer_schema)

Relation Instance

- The current values (*relation instance*) of a relation are specified by a table
- An element t of r is a *tuple*, represented by a *row* in a table

The diagram shows a table representing a relation instance. The table has three columns and four rows. Arrows point from the text 'attributes (or columns)' to the column headers. Arrows point from the text 'tuples (or rows)' to the rows. The table is labeled 'customer' below it.

<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
<i>Jones</i>	<i>Main</i>	<i>Harrison</i>
<i>Smith</i>	<i>North</i>	<i>Rye</i>
<i>Curry</i>	<i>North</i>	<i>Rye</i>
<i>Lindsay</i>	<i>Park</i>	<i>Pittsfield</i>

customer

Database

- A database consists of multiple relations
- Information about an enterprise is broken up into parts, with each relation storing one part of the information
 - account* : stores information about accounts
 - depositor* : stores information about which customer owns which account
 - customer* : stores information about customers
- Storing all information as a single relation such as *bank(account_number, balance, customer_name, ..)* results in repetition of information (e.g., two customers own an account) and the need for null values (e.g., represent a customer without an account)

Banking Example

branch (branch-name, branch-city, assets)

*customer (customer-name, customer-street,
customer-city)*

account (account-number, branch-name, balance)

loan (loan-number, branch-name, amount)

depositor (customer-name, account-number)

borrower (customer-name, loan-number)

Relational Algebra

- Primitives
 - Projection (π)
 - Selection (σ)
 - Cartesian product (\times)
 - Set union (\cup)
 - Set difference ($-$)
 - Rename (ρ)
- Other operations
 - Join (\bowtie)
 - Group by... aggregation
 - ...

What happens next?

- SQL
- System R (DB2), INGRES, ORACLE, SQL-Server, Teradata
 - B+-Tree (select)
 - Transaction Management
 - Join algorithm

Early 90's: OLAP & Data Warehouse

Database Workloads

- OLTP (online transaction processing)
 - Typical applications: e-commerce, banking, airline reservations
 - User facing: real-time, low latency, highly-concurrent
 - Tasks: relatively small set of “standard” transactional queries
 - Data access pattern: random reads, updates, writes (involving relatively small amounts of data)
- OLAP (online analytical processing)
 - Typical applications: business intelligence, data mining
 - Back-end processing: batch workloads, less concurrency
 - Tasks: complex analytical queries, often ad hoc
 - Data access pattern: table scans, large amounts of data involved per query

OLTP

- Most database operations involve *On-Line Transaction Processing* (OLTP).
 - Short, simple, frequent queries and/or modifications, each involving a small number of tuples.
 - Examples: Answering queries from a Web interface, sales at cash registers, selling airline tickets.

OLAP

- Of increasing importance are *On-Line Application Processing* (OLAP) queries.
 - Few, but complex queries --- may run for hours.
 - Queries do not depend on having an absolutely up-to-date database.


OLAP Examples

1. Amazon analyzes purchases by its customers to come up with an individual screen with products of likely interest to the customer.
2. Analysts at Wal-Mart look for items with increasing sales in some region.

One Database or Two?

- Downsides of co-existing OLTP and OLAP workloads
 - Poor memory management
 - Conflicting data access patterns
 - Variable latency
- Solution: separate databases
 - User-facing OLTP database for high-volume transactions
 - Data warehouse for OLAP workloads
 - How do we connect the two?

OLTP/OLAP Architecture




OLTP/OLAP Integration

- OLTP database for user-facing transactions
 - Retain records of all activity
 - Periodic ETL (e.g., nightly)
- Extract-Transform-Load (ETL)
 - Extract records from source
 - Transform: clean data, check integrity, aggregate, etc.
 - Load into OLAP database
- OLAP database for data warehousing
 - Business intelligence: reporting, ad hoc queries, data mining, etc.
 - Feedback to improve OLTP services

The Data Warehouse

- The most common form of data integration.
 - Copy sources into a single DB (*warehouse*) and try to keep it up-to-date.
 - Usual method: periodic reconstruction of the warehouse, perhaps overnight.
 - Frequently essential for analytic queries.

Warehouse Architecture



Star Schemas

- A *star schema* is a common organization for data at a warehouse. It consists of:
 1. *Fact table* : a very large accumulation of facts such as sales.
 - Often “insert-only.”
 2. *Dimension tables* : smaller, generally static information about the entities involved in the facts.

Example: Star Schema

- Suppose we want to record in a warehouse information about every beer sale: the bar, the brand of beer, the drinker who bought the beer, the day, the time, and the price charged.
- The fact table is a relation:
`Sales(bar, beer, drinker, day, time, price)`

Example, Continued

- The dimension tables include information about the bar, beer, and drinker “dimensions”:

Bars(bar, addr, license)

Beers(beer, manf)

Drinkers(drinker, addr, phone)

Visualization – Star Schema

Dimension Table (**Bars**)

--	--	--	--

Dimension Table (**Drinkers**)

--	--	--	--

Dimension Attrs.

Dependent Attrs.

--	--	--	--	--	--

Fact Table - **Sales**

--	--	--	--

Dimension Table (**Beers**)

--	--	--	--

Dimension Table (etc.)

Dimensions and Dependent Attributes

- Two classes of fact-table attributes:
 1. *Dimension attributes* : the key of a dimension table.
 2. *Dependent attributes* : a value determined by the dimension attributes of the tuple